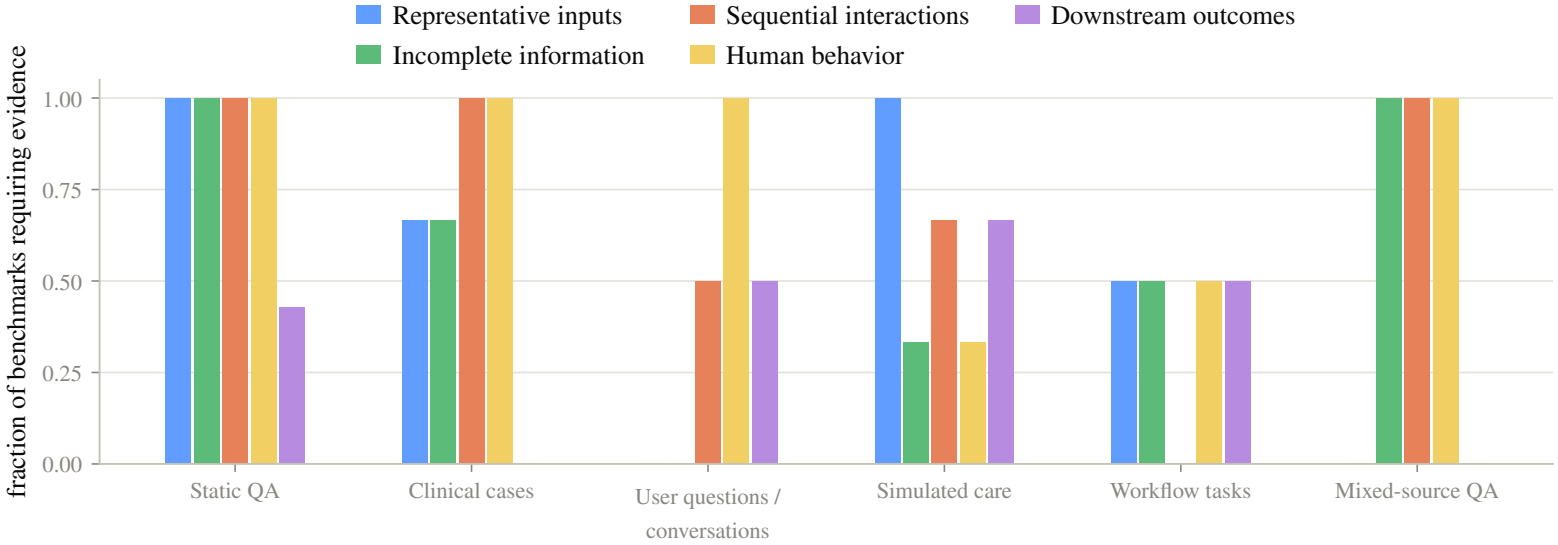
Representative inputs
Sequential interactions
Downstream outcomes
Incomplete information
Human behavior
fraction of benchmarks requiring evidence
1.00
0.75
0.50
0.25
0.00
Static QA
Clinical cases
User questions / conversations
Simulated care
Workflow tasks
Mixed-source QA